

Course Release Tools

Purpose

The course release tools separate authoring from publication. A working course repository can contain chapter source, speaker notes, slide source, code that has not yet been released, and other instructor material. The active course is the student-facing subset. It contains only the chapters and supporting files that have been selected for release.

This distinction matters for two reasons. First, the working repository remains the complete source of truth. Second, publishing a new chapter is an explicit operation. Material does not become public merely because it exists in the working repository.

The tools support this model without tying the build logic to one class. Course names, chapter lists, repository addresses, colors, file exclusions, and legacy URLs live in each course's `release.json`. The shared programs in this repository interpret that configuration and produce the site.

The release model

Each class uses three related locations:

1. The working repository contains all authoring material and all course code.
2. The working repository's `release/` directory is a generated preview of the student site.
3. The active repository contains the published subset and is served by GitHub Pages from its root.

The working repository is authoritative. Files in the active repository should be exact copies of their counterparts in `release/`. The intended difference is which chapters and files are present, not different versions of the same file.

The normal flow is:

```
working sources -> generated release -> local review -> active repository -> GitHub Pages
```

The tools currently build and check the first two steps. Copying the reviewed release into the active repository remains a deliberate operation so the result can be inspected before it becomes public.

Repository organization

The shared tools live in `~/courses/tools`:

```
tools/  
  build_course_release.py  
  check_course_release.py  
  code_viewer.html  
  md-to-pdf/  
  COURSE_RELEASE_TOOLS.md
```

Each working course repository contains its own `release.json`. It may also keep small `tools/build_release.py` and `tools/check_release.py` launchers so the same commands work from inside every class.

This arrangement puts shared behavior in one repository while leaving course policy with the course. A change to link checking or code display is made once. A change to the Advanced Database chapter list does not affect Structure of Programming Languages.

Building a course release

From anywhere, run the shared builder with the working repository path:

```
~/courses/tools/build_course_release.py ~/courses/structures/working
```

From a configured course repository, the compatibility launcher is shorter:

```
python3 tools/build_release.py
```

The builder reads `release.json`, runs any requested chapter or slide builds, and assembles the selected files. It creates chapter landing pages and code browsers, copies readings and slide assets, and writes the top-level course index. It builds into a temporary sibling directory first. The existing `release/` directory is replaced only after the new build completes.

For a quick preview that uses already-built PDFs and slides:

```
python3 tools/build_release.py --skip-build --output /tmp/course-preview
```

This is useful while changing site layout or release configuration. It avoids rebuilding documents whose source has not changed and leaves the normal release tree alone.

Checking a release

Run the checker after every build:

```
python3 tools/check_release.py
```

For a preview in another directory:

```
python3 tools/check_release.py --site /tmp/course-preview
```

The checker verifies that local links resolve, expected chapter directories are present, private source formats have not entered the site, released slide HTML does not contain speaker notes, and copied code files match their working-course sources byte for byte.

Application templates may contain routes such as `/pets` or template expressions that are meaningful only when the application runs. Those files are code, not site navigation, so the checker does not interpret links inside template directories as release-site links.

Configuring a course

The `release.json` file has two jobs. It describes the site, and it identifies the material to release.

The `site` object contains the course introduction, repository links, code-card description, and theme colors. Course-specific details remain here rather than inside the shared Python program.

The `chapters` array is the release boundary. Each entry may provide:

- a chapter PDF and an optional command that builds it;
- a slide deck, its build command, and supporting assets;
- a code directory;
- readings, including Markdown readings that must first become PDFs.

Removing a chapter from this list removes it from the generated release. Adding one makes it eligible for publication. This is why the configuration should be reviewed as carefully as the generated site.

The `release` object defines names and suffixes omitted while copying code. It is used to exclude authoring files, caches, local environments, database journals, and similar material. `legacy_urls` can preserve previously published workbook, slide, or topic paths when a course changes organization.

Markdown to PDF

The `md-to-pdf` directory contains the document publishing pipeline used by the course builders. The command is available at:

```
~/local/bin/md-to-pdf chapter.md
```

By default it writes `chapter.pdf` beside the Markdown source. An explicit output and an additional print stylesheet are also supported:

```
md-to-pdf chapter.md -o preview/chapter.pdf
md-to-pdf chapter.md --css course-print.css
```

Pandoc parses the Markdown into a document tree. The builder resolves local images relative to the source file and identifies Mermaid code blocks. Mermaid CLI renders those diagrams as Scalable Vector Graphics (SVG), preserving clean lines and text in the PDF. Pandoc then produces HTML, and WeasyPrint applies the shared print stylesheet and lays out Letter-sized pages.

The pipeline stages the completed PDF and replaces an existing output only after a successful render. A failed build therefore leaves the previous PDF intact. Missing local images and rendering errors stop the build instead of silently producing an incomplete document.

The source, lock files, templates, and styles are tracked in Git. Installed Node modules and the Python virtual environment live in the same directory for local use but are excluded from the repository. The paths under `~/ .local` are compatibility symlinks, so existing chapter build scripts do not need to change.

Publishing safely

A normal release should follow this sequence:

1. Build the working course's release.
2. Run the release checker.
3. Open the generated `index.html` locally and inspect the affected chapters.
4. Check PDFs and served slide decks, not only their source files.
5. Copy the reviewed subset into the active repository.
6. Confirm that active files match the generated release.
7. Commit and push the active repository.

The build and validation steps are intentionally separate. A successful build means the requested files could be assembled. A successful check means the assembled result satisfies the release rules that can be tested automatically. Neither replaces visual review of documents, slides, and pages.

Extending the tools

Shared behavior belongs in this repository. Course-specific wording and choices belong in `release.json`. That boundary keeps the tools reusable without making the configuration into a programming language of its own.

When a new course needs the same release model, start with a `release.json`, add the small launchers, and run a preview build outside the repository. Add a new configuration field only when the classes genuinely need different behavior. If all courses need the same behavior, change the shared implementation instead.